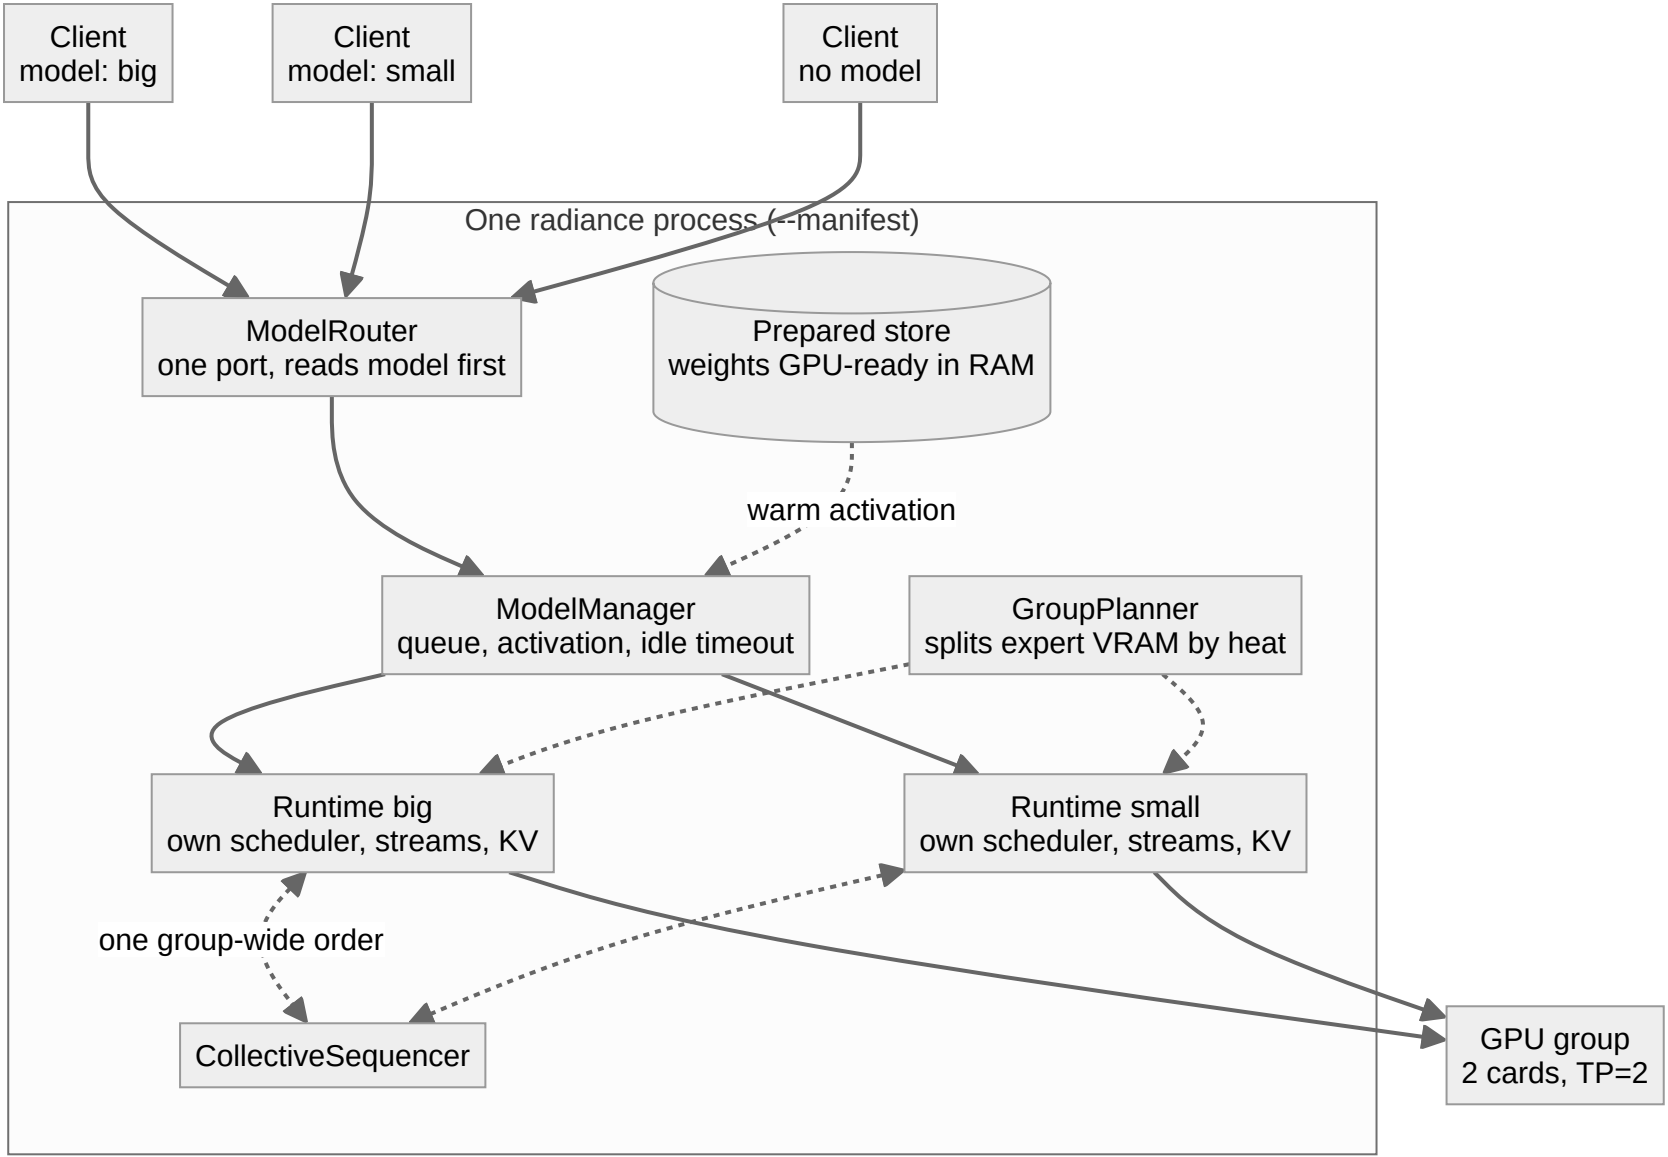
Client
model: big
Client
model: small
Client
no model
One radiance process (--manifest)
ModelRouter
one port, reads model first
Prepared store
weights GPU-ready in RAM
warm activation
ModelManager
queue, activation, idle timeout
GroupPlanner
splits expert VRAM by heat
Runtime big
own scheduler, streams, KV
Runtime small
own scheduler, streams, KV
one group-wide order
CollectiveSequencer
GPU group
2 cards, TP=2