Two models busy at once: who gets how much?

Dense 0.8B, one request per model (C1), tok/s · bar = sum of both models

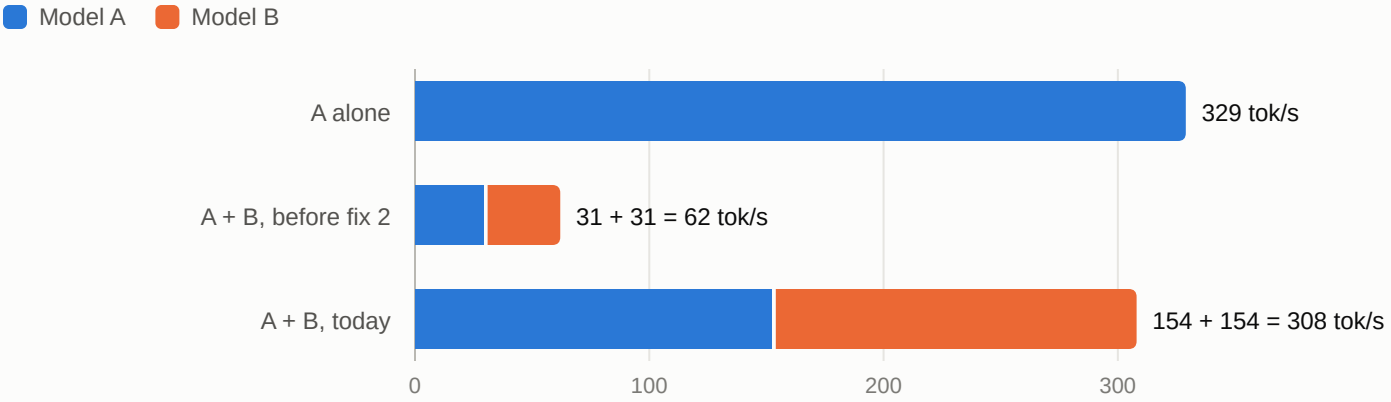

All models' collectives run in one group-wide order (ADR 0003): no deadlock, but no parallel communication either. Fix 2: the hand-off spins 300 µs before the thread sleeps (before: ~0.5 ms wake-up cost × 48 per decode step).