Elastic expert slab: VRAM follows the load

Two Flash-Next entries, rank 0, committed MiB per model · shared pool 20,680 MiB

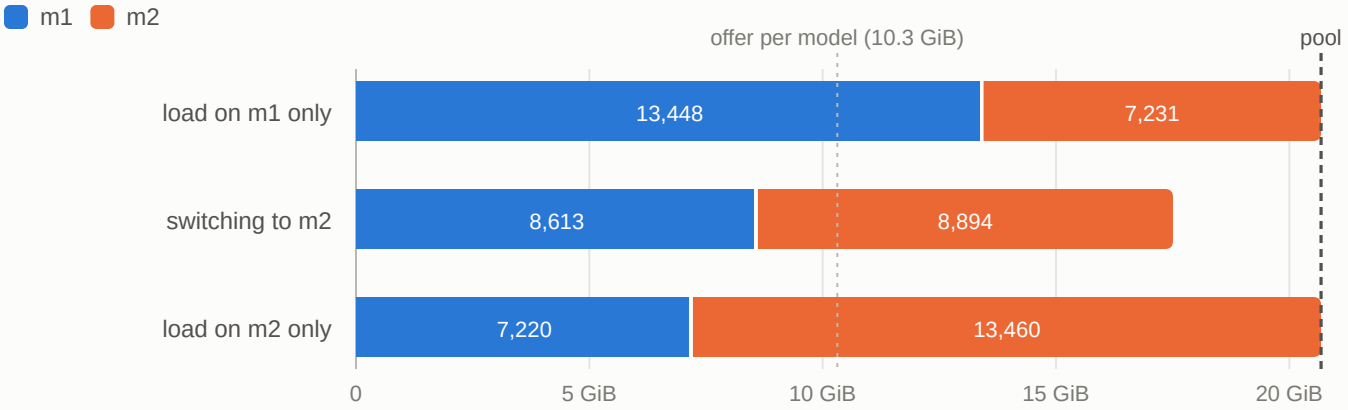

The busy model grows ~3 GiB past its offer, the idle one falls to its minimum. Σ committed ≤ pool in every sample. Moves in 128 MiB chunks, target reached in 20–30 s; greedy output of both models unchanged.