Where a model's weights can live

Faster at the top, larger at the bottom · bandwidths measured on 2× R9700 / PCIe

VRAM (31.9 GiB per card) — 600–870 GB/s

active models: static weights, KV cache, hot experts

Pinned host pool (RAM) — ≈ 28 GB/s over PCIe

RAM streaming: cold experts, static weights that do not fit in VRAM

Prepared store (RAM) — upload ≈ 0.25 s / 2.4 GB

deactivated models, packed GPU-ready: no parsing, no relayout

Disk / page cache — 2.5–3.1 GB/s measured

cold models: .rad container or HF checkpoint, full load path

States: active = VRAM (+ host pool) · prepared = store · cold = disk only. Page cache does not count as RAM backing.