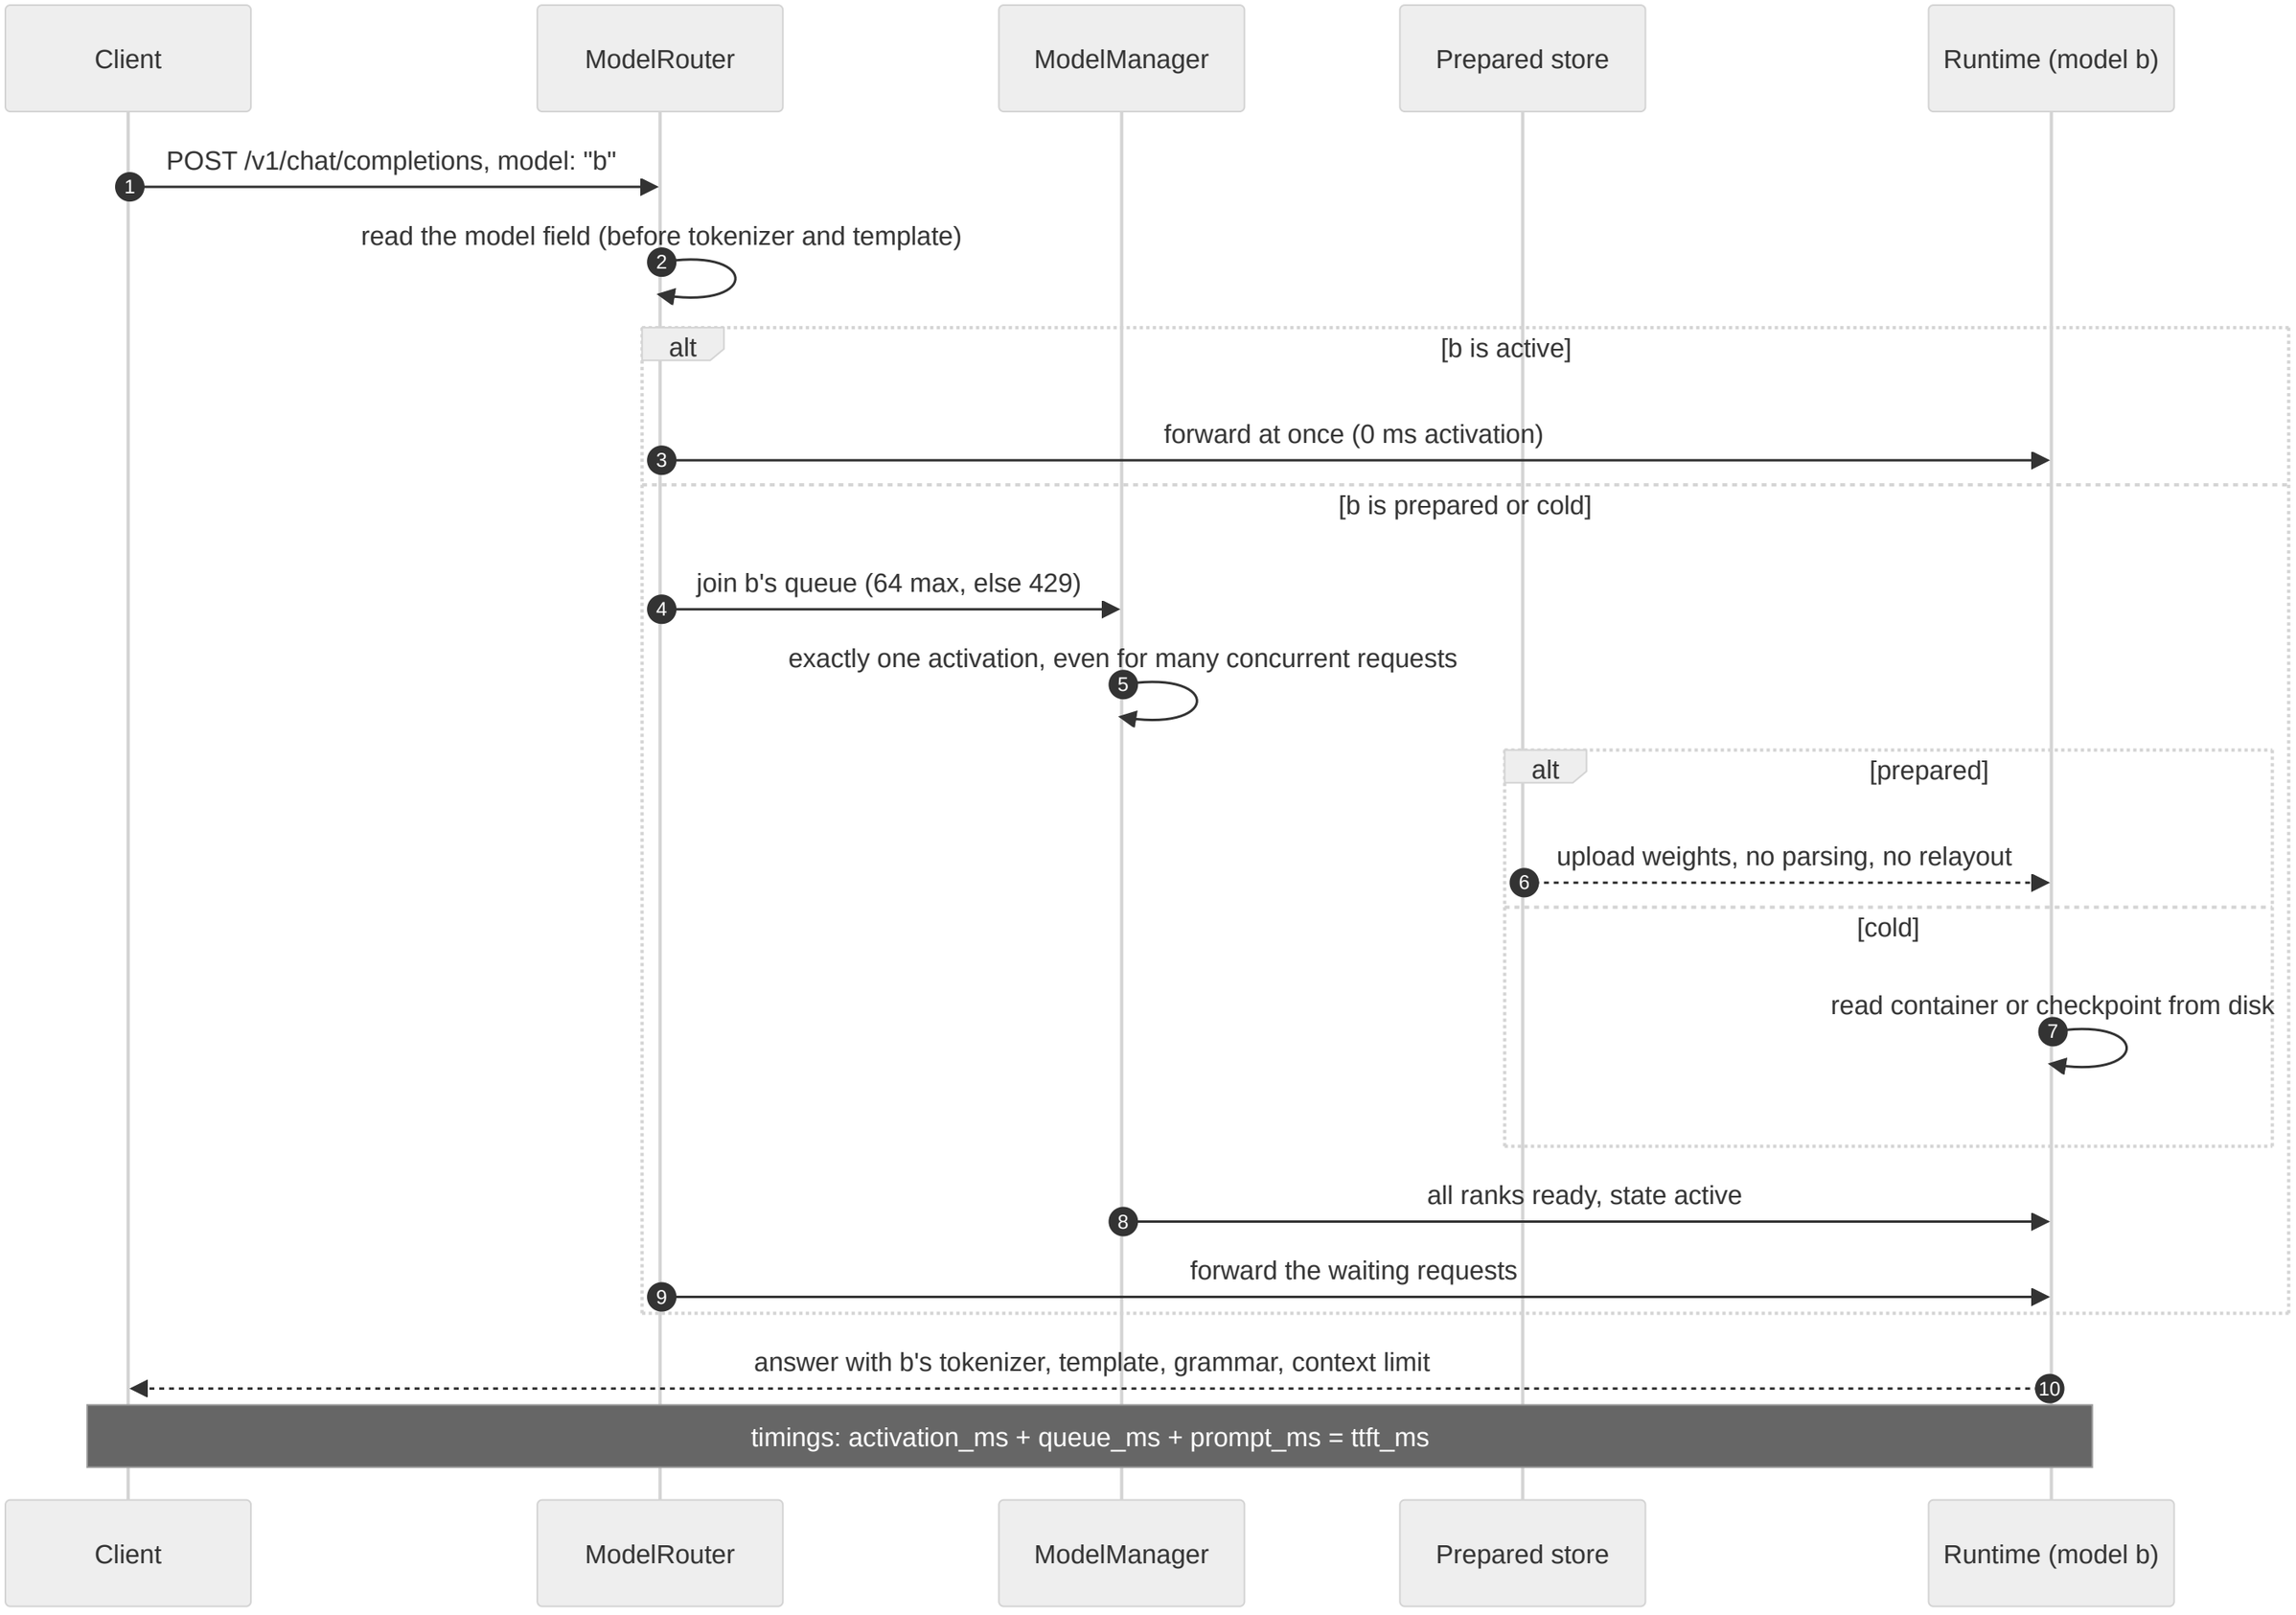
Client
ModelRouter
ModelManager
Prepared store
Runtime (model b)
POST /v1/chat/completions, model: "b"
1
read the model field (before tokenizer and template)
2
alt
[b is active]
forward at once (0 ms activation)
3
[b is prepared or cold]
join b's queue (64 max, else 429)
4
exactly one activation, even for many concurrent requests
5
alt
[prepared]
upload weights, no parsing, no relayout
6
[cold]
read container or checkpoint from disk
7
all ranks ready, state active
8
forward the waiting requests
9
answer with b's tokenizer, template, grammar, context limit
10
timings: activation_ms + queue_ms + prompt_ms = ttft_ms
Client
ModelRouter
ModelManager
Prepared store
Runtime (model b)