One card, two active models: how VRAM is shared

Schematic, not to scale

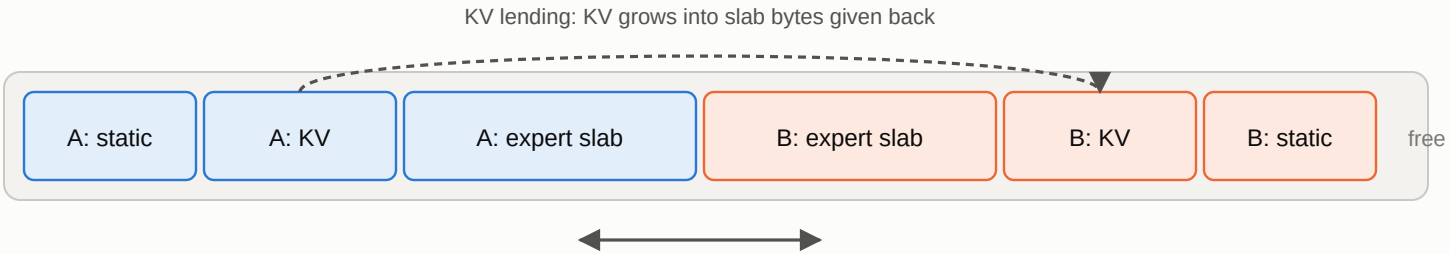

Order in which freed bytes are handed out

1 Streamed statics
back onto the card

→

2 KV cache
of a model with waiting prompts

→

3 Expert slab
growth by group heat

Running sequences never lose their KV cache; every model keeps a minimum share of experts.