Two ways to “one endpoint, many models”

llama-swap

Proxy in front of separate server processes

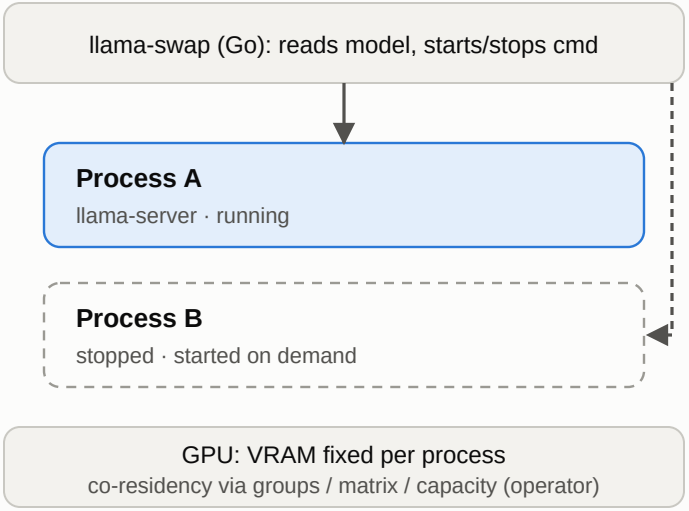

RadianceSwap

One process, one engine, shared cards

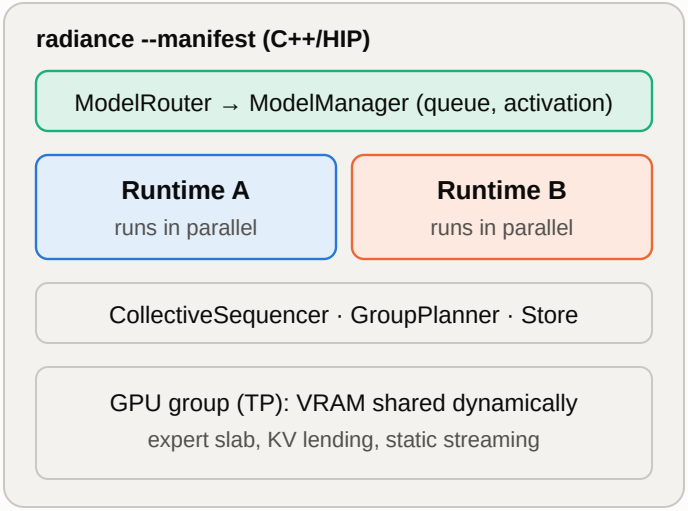

llama-swap: any OpenAI-compatible backend, any platform, mature. Switch = stop one process, start another.

RadianceSwap: Radiance engine only (ROCm), experimental. Switch = runtime in the same process, weights warm in RAM.